

# Lecture 5

---

## Last

- **Forwarding**: data plane
    - Directing one data packet
    - Each router using local routing state
  - **Routing**: control plane
    - Computing the forwarding tables that guide packets
    - Jointly computed by routers using a distributed protocol
  - Two correctness conditions for routing state:
    - no deadends
    - no loops
- 

## Goal (for routing)

- v1: Find a path to a given destination
- v2: Find a **least cost path** to a given destination

## Link costs, could represent

- propagation delay
- load
- cost

## **Least-cost path routing**

- Given: router graph & link costs
- Goal: find least-cost path
  - from each source router
  - to each destination router

### **“Least Cost” Routes**

- “Least cost” routes an easy way to avoid loops
    - No sensible cost metric is minimized by traversing the loop
  - Least cost routes are also “destination-based”
    - i.e., do not depend on the source
  - Least-cost paths form a spanning tree
- 

## **Approach 1: Link-state routing**

### **Routing algorithm for source u**

- Input: router graph & link costs
- Output: least cost path
  - from source router u
  - to every other router

### **Dijkstra’s algorithm**

- Source router considers every other router
  - starting from next “closest” neighbor

- Checks whether it can improve paths
  - by using that router as an intermediate point
- Ends when all intermediaries have been considered

## **From routing algorithm to protocol**

- Note: Dijkstra's is a **local** computation
  - computed by one node given complete network graph
- Possibilities
  - Option 1: a separate machine runs the algorithm
  - Option 2: every router runs the algorithm
- The internet currently uses Option 2

## **Link State Routing**

- Every router knows its local "link state"
  - router u: "(u,v) with cost=2; (u,x) with cost=1"
- A router floods its link state to all other routers
  - does so periodically or when its link state changes
- Hence, every router learns the entire network graph
  - runs Dijkstra's locally to compute forwarding table
- OSPF is a link-state protocol (IETF RFC 2328 or 5340)
  - Berkeley runs OSPF internally!

## **Convergence**

- All routers have consistent routing information
  - E.g., all nodes having the same link-state database
- Forwarding is consistent after convergence

- All nodes have the same link-state database
- All nodes forward packets on same paths

## **Convergence Delay**

- Time to achieve convergence
- Sources of convergence delay?
  - Time to detect failure
  - Time to flood link-state information
  - Time to re-compute forwarding tables
- Performance during convergence period?
  - Lost packets due to blackholes
  - Looping packets
  - Out-of-order packets reaching the destination

## **Link State Routing**

- Are loops possible?
    - yes, until convergence
  - Scalability
    - Messages:  $O(N^2)$
    - Computation time:  $O(N^2)$
    - Convergence time:  $O(\text{network diameter})$
    - Entries in forwarding table:  $O(N)$
  - Do we have to use Dijkstra's? No.
-

## Approach 2: Distance-vector routing

### Experiment

- Your job: find the youngest person in the room
- Ground rules
  - You may not leave your seat, nor shout loudly across the class
  - You may talk with your immediate neighbors (hint: “exchange updates” with them)
- At the end of **5 minutes**, I will pick a victim and ask:
  - Who is the youngest person in the room? (name, date)
  - Which one of your neighbors first told you this information?

### Distance-vector routing

- Input to each router
  - local link costs & neighbor messages
- Output of each router
  - least-cost path to every other router
- Distributed algorithm
- All routers run it “together”
  - each router runs its own instance
  - neighbors exchange and react to each other’s message

### Bellman-Ford algorithm

- All neighbors exchange information
  - Each router checks whether it can improve current paths by leveraging the new information
- Ends when no improvement is possible

## **Bellman-Ford equation**

- $d_x(z) = \min_n \{ \text{cost}(x,n) + d_n(z) \}$ 
  - for all neighbors  $n$
- Formalizes the following decision:
  - Pick as the next hop for destination  $z$  the neighbor that results in the least-cost path  $z$

## **Problems with Bellman-Ford**

- Routing loops
  - $z$  routes through  $y$ ,  $y$  routes through  $x$
  - $y$  loses connectivity to  $x$
  - $y$  decides to route through  $z$
- Can take a very long time to resolve
  - Count-to-infinity scenario

## **Solution**

- Poisoned reverse
  - if  $z$  routes to  $x$  through  $y$ ,  $z$  advertises to  $y$  that its cost to  $x$  is infinite
  - $y$  never decides to route to  $x$  through  $z$
- Often avoids the count-to-infinity problem

## **Distance Vector Routing**

- Are loops possible?
  - Yes, until convergence
  - Convergence slower than in Link-State

- Scalability
  - Requires fewer messages than Link-State
  - Update time on arrival of a new DV from neighbor:  $O(N)$
  - Convergence time:  $O(\text{network diameter})$
  - Entries in forwarding table:  $O(N)$
- Rip is a protocol that implements DV (IETF RFC 2080)